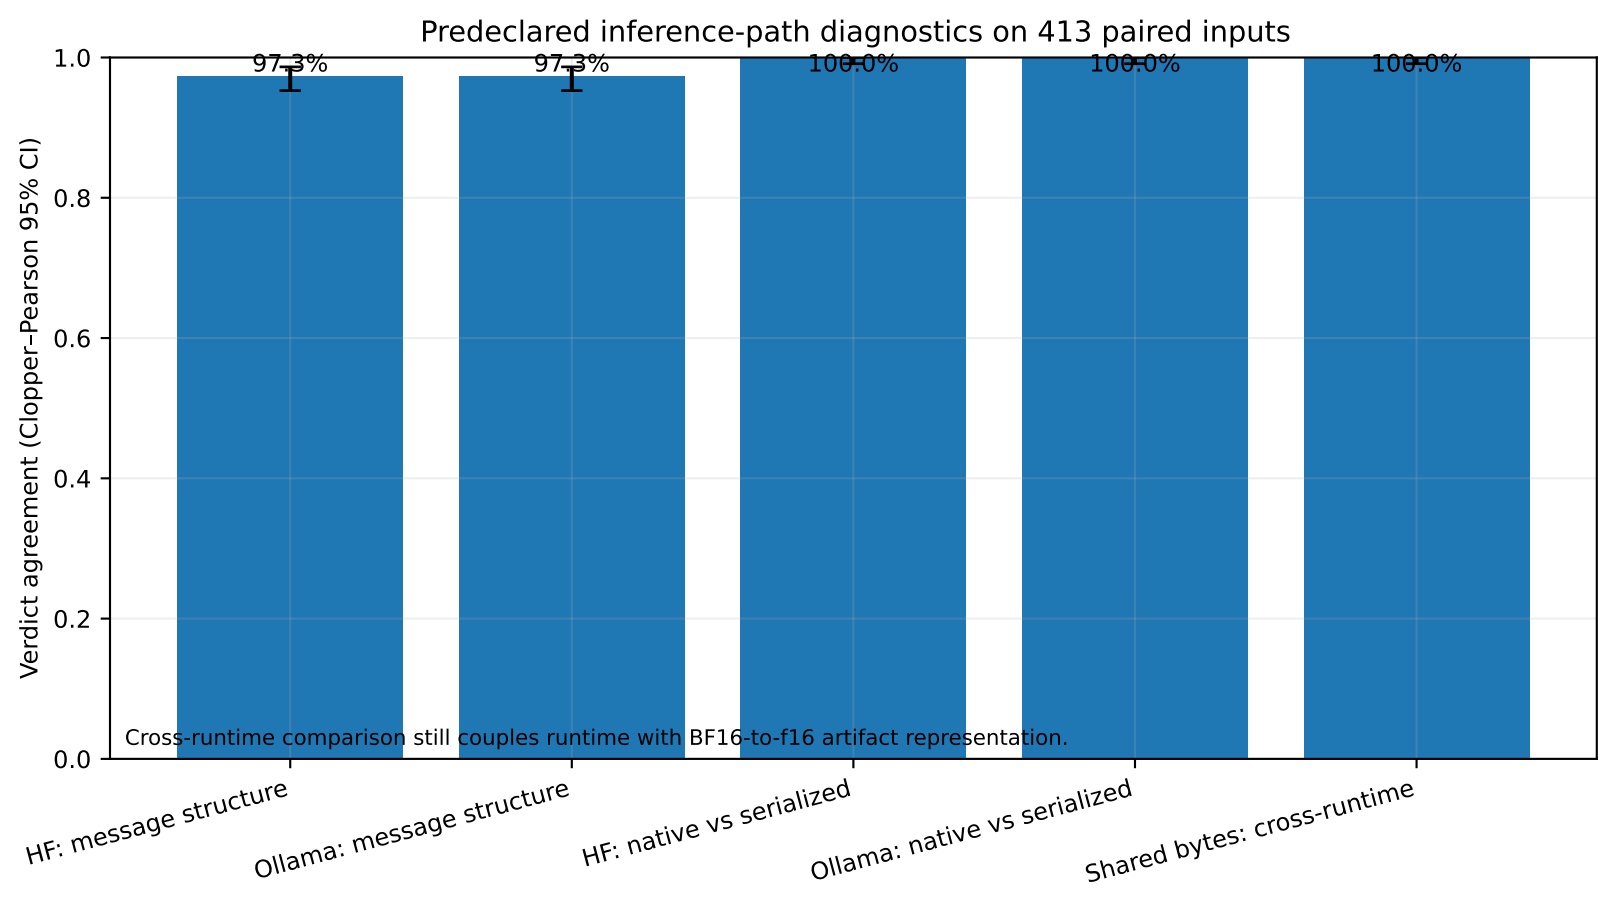
Predeclared inference-path diagnostics on 413 paired inputs
Verdict agreement (Clopper–Pearson 95% CI)
1.0
0.8
0.6
0.4
0.2
0.0
97.3%
97.3%
100.0%
100.0%
100.0%
Cross-runtime comparison still couples runtime with BF16-to-f16 artifact representation.
HF: message structure
Ollama: message structure
HF: native vs serialized
Ollama: native vs serialized
Shared bytes: cross-runtime